



# Optimization

*Minjong Lee*


*Pohang University of Science and Technology (Postech) CSE*

[Minjong.Lee@postech.ac.kr](mailto:Minjong.Lee@postech.ac.kr)

# Training/Test error and Generalization


- **Training error**
  - error of training dataset
- **Test error**
  - Error of test dataset (new and previously unseen data)
- **Generalization**
  - The ability to perform well on previously unobserved input

# Training/Test error and Generalization

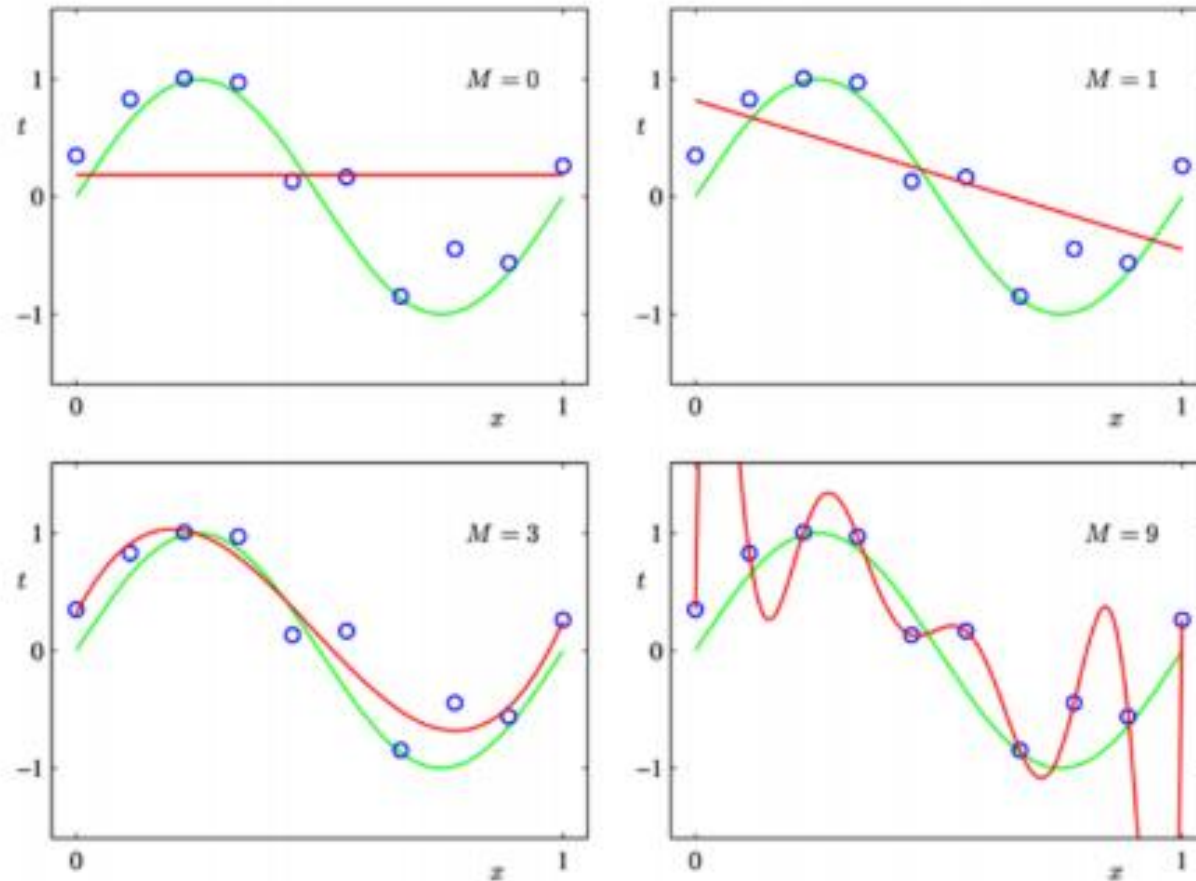


# Overfitting and Underfitting


- Underfitting
  - The training data and test data have high error rates.
- Overfitting
  - The training data has a low error rate but the test data has a high error rate.



# Overfitting and Underfitting



# Overfitting and Underfitting




# Contents

- Activation function
- Data preprocessing
- Batch Normalization
- Weight initialization
- Fancy optimizers

# Activation function


## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$




## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$




## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

## ELU



$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$






# Activation function: Sigmoid

- $\sigma(x) = 1/(1 + e^{-x})$
- range [0, 1]
- Problem
  - Gradient Vanishing
  - Exp() is very expensive.




# Activation function: ReLU

- $f(x) = \max(0, x)$
- Does not saturate (in + region)
- Very computationally efficient
- Problem
  - Dead ReLU




# Activation function: ReLU


- Dead ReLU



# Activation function: LeakyReLU



*ReLU activation function*



*LeakyReLU activation function*

# Activation functions

Identity	Sigmoid	TanH	ArcTan
ReLU	Leaky ReLU	Randomized ReLU	Parameteric ReLU
Binary	Exponential Linear Unit	Soft Sign	Inverse Square Root Unit (ISRU)
Inverse Square Root Linear	Square Non-Linearity	Bipolar ReLU	Soft Plus


# Contents

- Activation function
- Data preprocessing
- Batch Normalization
- Weight initialization
- Fancy optimizers

# Data problem: Scale

- data
  - (height(m), weight(kg))
  - ex) (1.5m, 70kg)
- The network will be biased to the 'weight'.
- To avoid this problem,
  - we must normalize the data.

# Data preprocessing




Assume X is data matrix, each sample in a row




# Data preprocessing

**Before normalization:** classification loss very sensitive to changes in weight matrix; hard to optimize



**After normalization:** less sensitive to small changes in weights; easier to optimize






# Contents

- Activation function
- Data preprocessing
- Batch Normalization
- Weight initialization
- Fancy optimizers

# Covariate shift

- There are unexpected shifts in the distribution of layers' input.



# Batch normalization

- Input:  $X \in \mathbb{R}^{N \times D}$
- What if zero-mean, unit var is too hard of a constraints?
- Learnable scale and shift parameters:  
$$\gamma, \beta \in \mathbb{R}^D$$
- Learning  $\gamma = \sigma, \beta = \mu$  will recover the identity function!

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{ij}$$

Per-feature mean.  
Shape is  $D$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{ij} - \mu_j)^2$$

Per-feature var.  
Shape is  $D$

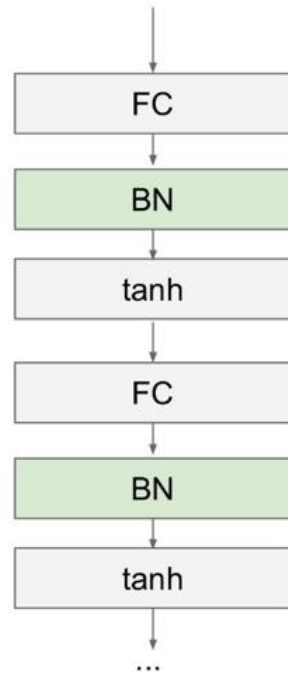
$$\hat{x}_{ij} = \frac{x_{ij} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

Normalized x.  
Shape is  $N \times D$

$$y_{ij} = \gamma_j \hat{x}_{ij} + \beta_j$$





Output.  
Shape is  $N \times D$

# Batch normalization



- Makes deep networks much easier to train!
- Improves gradient flow
- Allows higher learning rates, faster convergence
- Networks become more robust to initialization
- Acts as regularization during training
- Zero overhead at test-time: can be fused with conv!
- Behaves differently during training and testing:  
this  
is a very common source of bugs!

# Other normalization techniques




# Contents

- Activation function
- Data preprocessing
- Batch Normalization
- Weight initialization
- Fancy optimizers


# Weight initialization

- Init weights with  $\mathcal{N}(0, 0.01)$  with **tanh**(x),

**tanh**  
 $\tanh(x)$



- All zero, no learning!







# Weight initialization

- Init weights with  $\mathcal{N}(\mathbf{0}, \mathbf{1})$  with **tanh(x)**,
- Local gradient all zero, no learning!

**tanh**  
 $\tanh(x)$




$$\frac{d}{dx} \tanh(x) = 1 - \tanh^2(x)$$




# Xavier initialization

- Init weights with  $\mathcal{N}(0, \text{Var}(W))$  with  $\tanh(x)$ ,
  - $\text{Var}(W) = \frac{1}{\sqrt{D_{in}}}$
  - $D_{in}$  : the size of the dimension of input




# Xavier initialization

- Init weights with  $\mathcal{N}(\mathbf{o}, \text{Var}(\mathbf{W}))$  with  $\tanh(\mathbf{x})$ ,
  - $\text{Var}(\mathbf{W}) = \frac{1}{\sqrt{D_{in}}}$
  - $D_{in}$  : the size of the dimension of input



# He initialization



- Init weights with  $\mathcal{N}(\mathbf{o}, \text{Var}(\mathbf{W}))$  with  $\tanh(\mathbf{x})$ ,
  - $\text{Var}(\mathbf{W}) = \frac{2}{\sqrt{D_{in}}}$
  - $D_{in}$  : the size of the dimension of input



# Contents

- Activation function
- Data preprocessing
- Batch Normalization
- Weight initialization
- Fancy optimizers


# Stochastic Gradient Descent (SGD)



$$\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial L}{\partial \mathbf{W}}$$


# Stochastic Gradient Descent (SGD)

- Problem



# Problem of gradient descent



- Zero gradient, gradient descent gets stuck.








# Momentum




# SGD + Momentum

- SGD with momentum is better than only SGD.




$$V_t = \beta V_{t-1} + \alpha \nabla_w L(W, X, y)$$
$$W = W - V_t$$


# Importance of learning rate




Big Learning Rate



Just right



Too small



# Fancy optimizer: AdaGrad (Adaptive Gradients)


- Learning rate decay
  - Decay the learning rate for parameters in proportion to their update history (more updates means more decay).
- When  $h$  is very big, the parameter will start receiving very small updates.

$$h \leftarrow h + \frac{\partial L}{\partial W} \odot \frac{\partial L}{\partial W}$$
$$W \leftarrow W - \eta \frac{1}{\sqrt{h}} \frac{\partial L}{\partial W}$$

# Fancy optimizer: RMSProp

- Everything is very similar to AdaGrad, except now we decay the denominator as well.

$$\begin{aligned}h &\leftarrow h + \frac{\partial L}{\partial W} \odot \frac{\partial L}{\partial W} \\W &\leftarrow W - \eta \frac{1}{\sqrt{h}} \frac{\partial L}{\partial W}\end{aligned}$$



$$\begin{aligned}h_i &\leftarrow \rho h_{i-1} + (1 - \rho) \frac{\partial L_i}{\partial W} \odot \frac{\partial L_i}{\partial W} \\W &\leftarrow W - \eta \frac{1}{\sqrt{h}} \frac{\partial L}{\partial W}\end{aligned}$$

# Fancy optimizer: Adam (Adaptive moment)

- RMSProp(learning rate decay) + Momentum

## 알고리즘 5-5 Adam


입력: 훈련집합  $\mathbb{X}$ ,  $\mathbb{Y}$ , 학습률  $\rho$ , 모멘텀 계수  $\alpha_1$ , 가중 이동 평균 계수  $\alpha_2$

출력: 최적의 매개변수  $\hat{\Theta}$

```
1  난수를 생성하여 초기해  $\Theta$ 를 설정한다.
2   $\mathbf{v} = \mathbf{0}, \mathbf{r} = \mathbf{0}$ 
3   $t = 1$ 
4  repeat
5    그래디언트  $\mathbf{g} = \frac{\partial J}{\partial \Theta} \Big|_{\Theta}$ 를 구한다.
6     $\mathbf{v} = \alpha_1 \mathbf{v} - (1 - \alpha_1) \mathbf{g}$  // 속도 벡터
7     $\mathbf{v} = \frac{1}{1 - (\alpha_1)^t} \mathbf{v}$ 
8     $\mathbf{r} = \alpha_2 \mathbf{r} + (1 - \alpha_2) \mathbf{g} \odot \mathbf{g}$  // 그래디언트 누적 벡터
9     $\mathbf{r} = \frac{1}{1 - (\alpha_2)^t} \mathbf{r}$ 
10    $\Delta \Theta = -\frac{\rho}{\epsilon + \sqrt{\mathbf{r}}} \mathbf{v}$ 
11    $\Theta = \Theta + \Delta \Theta$ 
12    $t++$ 
13 until (멈춤 조건)
14  $\hat{\Theta} = \Theta$ 
```

# Scheduling learning rate

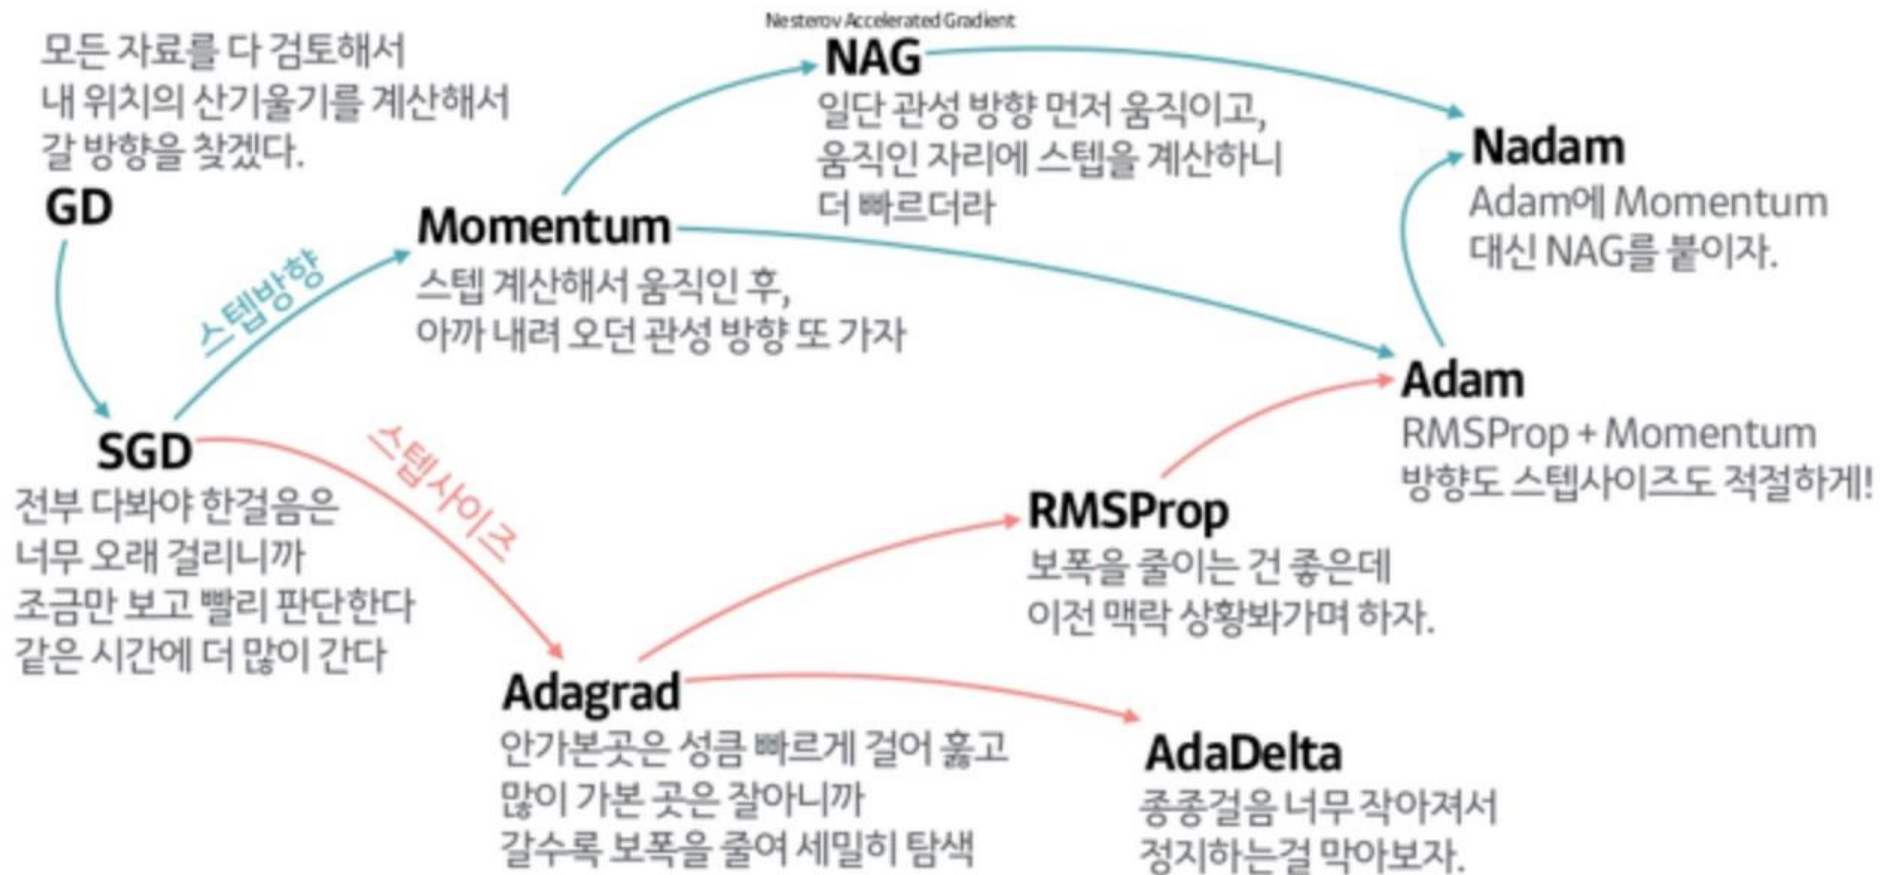
- SGD, SGD+Momentum, Adagrad, RMSProp, Adam all have learning rate as a hyperparameter.




When we train NN, we typically start with large learning rate and decay over time.



# Fancy optimizers



# Early Stopping



# Practice

- Activation
- Batch Normalization
- Weight Initialization
- Optimizers

# Activation

## Functions

`deserialize(...)` : Returns activation function given a string identifier.

`elu(...)` : Exponential Linear Unit.

`exponential(...)` : Exponential activation function.

`gelu(...)` : Applies the Gaussian error linear unit (GELU) activation function.

`get(...)` : Returns function.

`hard_sigmoid(...)` : Hard sigmoid activation function.

`linear(...)` : Linear activation function (pass-through).

`relu(...)` : Applies the rectified linear unit activation function.

`selu(...)` : Scaled Exponential Linear Unit (SELU).

`serialize(...)` : Returns the string identifier of an activation function.

`sigmoid(...)` : Sigmoid activation function,  $\text{sigmoid}(x) = 1 / (1 + \exp(-x))$ .

`softmax(...)` : Softmax converts a vector of values to a probability distribution.

`softplus(...)` : Softplus activation function,  $\text{softplus}(x) = \log(\exp(x) + 1)$ .

`softsign(...)` : Softsign activation function,  $\text{softsign}(x) = x / (\text{abs}(x) + 1)$ .

`swish(...)` : Swish activation function,  $\text{swish}(x) = x * \text{sigmoid}(x)$ .

`tanh(...)` : Hyperbolic tangent activation function.

[https://www.tensorflow.org/api\\_docs/python/tf/keras/activations](https://www.tensorflow.org/api_docs/python/tf/keras/activations)

# Batch Normalization

```
tf.keras.layers.BatchNormalization(  
    axis=-1,  
    momentum=0.99,  
    epsilon=0.001,  
    center=True,  
    scale=True,  
    beta_initializer='zeros',  
    gamma_initializer='ones',  
    moving_mean_initializer='zeros',  
    moving_variance_initializer='ones',  
    beta_regularizer=None,  
    gamma_regularizer=None,  
    beta_constraint=None,  
    gamma_constraint=None,  
    **kwargs  
)
```

[https://www.tensorflow.org/api\\_docs/python/tf/keras/layers/BatchNormalization](https://www.tensorflow.org/api_docs/python/tf/keras/layers/BatchNormalization)

# Weight Initialization

## Classes

`class Constant` : Initializer that generates tensors with constant values.

`class GlorotNormal` : The Glorot normal initializer, also called Xavier normal initializer.

`class GlorotUniform` : The Glorot uniform initializer, also called Xavier uniform initializer.

`class HeNormal` : He normal initializer.

`class HeUniform` : He uniform variance scaling initializer.

`class Identity` : Initializer that generates the identity matrix.

`class Initializer` : Initializer base class: all Keras initializers inherit from this class.

`class LecunNormal` : Lecun normal initializer.

`class LecunUniform` : Lecun uniform initializer.

`class Ones` : Initializer that generates tensors initialized to 1.

`class Orthogonal` : Initializer that generates an orthogonal matrix.

`class RandomNormal` : Initializer that generates tensors with a normal distribution.

`class RandomUniform` : Initializer that generates tensors with a uniform distribution.

`class TruncatedNormal` : Initializer that generates a truncated normal distribution.

`class VarianceScaling` : Initializer capable of adapting its scale to the shape of weights tensors.

`class Zeros` : Initializer that generates tensors initialized to 0.

`class constant` : Initializer that generates tensors with constant values.

[https://www.tensorflow.org/api\\_docs/python/tf/keras/initializers](https://www.tensorflow.org/api_docs/python/tf/keras/initializers)

# Weight Initialization

- Xavier initialization

```
fc = layers.Dense(128, kernel_initializer=tf.keras.initializers.GlorotNormal)
```



```
fc = layers.Dense(128, kernel_initializer='glorot_normal')
```

# Weight Initialization

- Normal distribution initialization
  - `layers.Dense(128, kernel_initializer='normal')`
- Xavier initialization
  - `layers.Dense(128, kernel_initializer='glorot_normal')`
- He initialization
  - `layers.Dense(128, kernel_initializer='he_normal')`



# Optimizers

## Classes

`class Adadelta` : Optimizer that implements the Adadelta algorithm.

`class Adagrad` : Optimizer that implements the Adagrad algorithm.

`class Adam` : Optimizer that implements the Adam algorithm.

`class Adamax` : Optimizer that implements the Adamax algorithm.

`class Ftrl` : Optimizer that implements the FTRL algorithm.

`class Nadam` : Optimizer that implements the NAdam algorithm.

`class Optimizer` : Base class for Keras optimizers.

`class RMSprop` : Optimizer that implements the RMSprop algorithm.

`class SGD` : Gradient descent (with momentum) optimizer.

[https://www.tensorflow.org/api\\_docs/python/tf/keras/optimizers](https://www.tensorflow.org/api_docs/python/tf/keras/optimizers)

# Practice

- With using skills that we learned, try to upgrade the network for better performance!